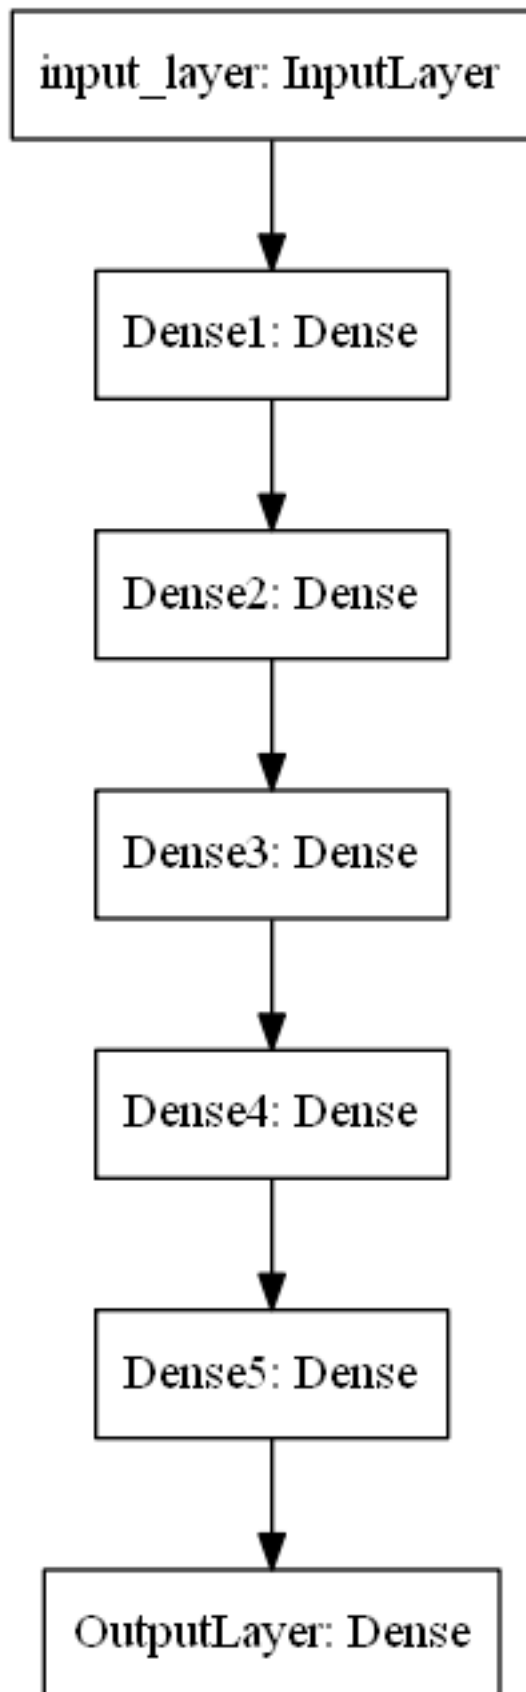


# TensorFlow-Callbacks

October 9, 2022

1. Download the data from [here](#). You have to use data.csv file for this assignment
2. Code the model to classify data like below image. You can use any number of units in your Dense layers.



### 3. Writing Callbacks

- Write your own callback function, that has to print the micro F1 score and AUC score after each epoch. Do not use `tf.keras.metrics` for calculating AUC and F1 score.
- Save your model at every epoch if your validation accuracy is improved from previous epoch.
- You have to decay learning based on below conditions
  - Cond1. If your validation accuracy at that epoch is less than previous epoch accuracy, you have to decrease the learning rate by 10%.
  - Cond2. For every 3rd epoch, decay your learning rate by 5%.
- If you are getting any NaN values(either weights or loss) while training, you have to terminate your training.
- You have to stop the training if your validation accuracy is not increased in last 2 epochs.
- Use tensorboard for every model and analyse your scalar plots and histograms (you need to upload the screenshots and write the observations for each model for evaluation).

```
[1]: import warnings
warnings.filterwarnings('ignore')
```

```
[2]: from google.colab import drive
drive.mount(mountpoint='/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call `drive.mount("/content/drive", force_remount=True)`.

```
[3]: from IPython.display import display
```

```
[4]: from matplotlib import pyplot as plt
from matplotlib import style
style.use(style='seaborn-deep')
```

```
[5]: from sklearn.metrics import auc, f1_score, roc_curve
from sklearn.model_selection import train_test_split
```

```
[6]: from tensorflow.keras.layers import Input, Dense
from tensorflow.keras.models import Model
from tensorflow.keras.callbacks import EarlyStopping, LearningRateScheduler, \
    ModelCheckpoint, ReduceLROnPlateau
```

```
[7]: from tqdm import tqdm
```

```
[8]: import datetime
import numpy as np
import os
import pandas as pd
import pickle
import seaborn as sns
```

```
import tensorflow as tf
```

```
[9]: data_file = '/content/drive/MyDrive/Applied-AI/Assignment-20/data.csv'
```

```
[10]: data_df = pd.read_csv(filepath_or_buffer=data_file)
display(data_df.head())
```

	f1	f2	label
0	0.450564	1.074305	0.0
1	0.085632	0.967682	0.0
2	0.117326	0.971521	1.0
3	0.982179	-0.380408	0.0
4	-0.720352	0.955850	0.0

```
[11]: display(data_df.describe())
```

	f1	f2	label
count	20000.000000	20000.000000	20000.000000
mean	0.000630	-0.000745	0.500000
std	0.671165	0.674704	0.500013
min	-1.649781	-1.600645	0.000000
25%	-0.589878	-0.596424	0.000000
50%	0.001795	-0.003113	0.500000
75%	0.586631	0.597803	1.000000
max	1.629722	1.584291	1.000000

```
[12]: display(data_df['label'].value_counts().to_frame())
```

label	
0.0	10000
1.0	10000

We have a balanced data.

```
[13]: X = data_df.drop(columns=['label'])
y = data_df['label'].values

display(X.head())
print(y[:5])
```

	f1	f2
0	0.450564	1.074305
1	0.085632	0.967682
2	0.117326	0.971521
3	0.982179	-0.380408
4	-0.720352	0.955850

[0. 0. 1. 0. 0.]

```
[14]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33,
↳random_state=0)
```

```
[15]: print(X_train.shape, y_train.shape)
print(X_test.shape, y_test.shape)
```

(13400, 2) (13400,)

(6600, 2) (6600,)

Custom callbacks.

```
[16]: class ScoreCallBack(tf.keras.callbacks.Callback):
    """
    This class is a custom callback for scores.
    """

    def __init__(self, validation_data):
        self.X_test = validation_data[0]
        self.y_test = validation_data[1]

    def on_epoch_end(self, epoch, logs={}):
        """
        This overrides the method present in parent class and prints F1 Score,
↳and AUC Score on each epoch end.
        """
        y_pred_probs = self.model.predict(x=self.X_test)
        y_pred = np.argmax(a=y_pred_probs, axis=1)

        fpr, tpr, thresholds = roc_curve(y_true=self.y_test, y_score=y_pred)
        auc_score_val = auc(x=fpr, y=tpr)
        f1_score_val = f1_score(y_true=self.y_test, y_pred=y_pred)

        print("\nAUC: {} - F1-Score: {}".format(auc_score_val, f1_score_val))
```

```
[17]: class TerminateCallBack(tf.keras.callbacks.Callback):
    """
    This class is a custom callback for termination if NaN.
    """

    def on_epoch_end(self, epoch, logs={}):
        """
        This overrides the method present in parent class and terminates if
↳loss is NaN.
        """
        loss = logs.get('loss')
        model_weights = self.model.get_weights()
        if ((loss is not None) or
            (model_weights is not None)):
```

```

        if (np.isnan(loss) or
            np.isinf(loss) or
            np.any([np.any(np.isnan(x)) for x in model_weights])):
            print("Invalid loss/weights and terminated at epoch {}".format(epoch))
            self.model.stop_training = True

```

```

[18]: def scheduler(epoch, lr):
        """
        This function decreases the learning rate for every 3rd epoch by 5%.
        """
        if (epoch + 1) % 3 == 0:
            lr -= (lr * 0.05)
            return lr
        else:
            return lr

```

Model-1

1. Use tanh as an activation for every layer except output layer.
2. Use SGD with momentum as optimizer.
3. Use RandomUniform(0,1) as initializer.
4. Analyze your output and training process.

```

[19]: %load_ext tensorboard
        # Clear any logs from previous runs
        !rm -rf ./logs/

```

```

[20]: m1_ki = tf.keras.initializers.RandomUniform(minval=0, maxval=1)

```

```

[21]: input_layer = Input(shape=(2,))
        layer1 = Dense(units=20, activation='tanh',
            ↪kernel_initializer=m1_ki)(input_layer)
        layer2 = Dense(units=15, activation='tanh', kernel_initializer=m1_ki)(layer1)
        layer3 = Dense(units=10, activation='tanh', kernel_initializer=m1_ki)(layer2)
        layer4 = Dense(units=5, activation='tanh', kernel_initializer=m1_ki)(layer3)
        layer5 = Dense(units=2, activation='tanh', kernel_initializer=m1_ki)(layer4)
        output_layer = Dense(units=1, activation='sigmoid',
            ↪kernel_initializer=m1_ki)(layer5)

```

```

[22]: model1 = Model(inputs=input_layer, outputs=output_layer)
        optimizer = tf.keras.optimizers.SGD(learning_rate=0.01, momentum=0.9)
        model1.compile(optimizer=optimizer, loss='binary_crossentropy',
            ↪metrics=['accuracy', 'AUC'])

```

```

[23]: score_callback = ScoreCallBack(validation_data=(X_test, y_test))

        termination_callback = TerminateCallBack()

```

```

decay_lr_callback = ReduceLROnPlateau(monitor='val_accuracy', factor=0.1,
    ↪patience=1, verbose=1, mode='auto')

lr_schedule_callback = LearningRateScheduler(schedule=scheduler, verbose=1)

filepath = "model_save/weights-{epoch:02d}-{val_accuracy:.4f}.hdf5"
model_save_callback = ModelCheckpoint(filepath=filepath,
    ↪monitor='val_accuracy', verbose=1, save_best_only=True, mode='auto')

early_stop_callback = EarlyStopping(monitor='val_accuracy', min_delta=0.01,
    ↪patience=2, verbose=1)

log_dir = os.path.join('logs', 'fits', datetime.datetime.now().
    ↪strftime("%Y%m%d-%H%M%S"))
tensorboard_callback = tf.keras.callbacks.TensorBoard(log_dir=log_dir,
    ↪histogram_freq=1)

callbacks = [score_callback, termination_callback,
             decay_lr_callback, lr_schedule_callback,
             model_save_callback, early_stop_callback, tensorboard_callback]

```

```

[24]: model1.fit(x=X_train, y=y_train, epochs=5, validation_data=(X_test, y_test),
    ↪batch_size=20, callbacks=callbacks)

```

Epoch 1: LearningRateScheduler setting learning rate to 0.009999999776482582.

Epoch 1/5

1/670 [...] - ETA: 18:40 - loss: 0.9227 - accuracy:  
0.5000 - auc: 0.4835

WARNING:tensorflow:Callback method `on\_train\_batch\_end` is slow compared to the batch time (batch time: 0.0044s vs `on\_train\_batch\_end` time: 0.0134s). Check your callbacks.

663/670 [=====>.] - ETA: 0s - loss: 0.6998 - accuracy:  
0.5020 - auc: 0.4996  
AUC: 0.5 - F1-Score: 0.0

Epoch 1: val\_accuracy improved from -inf to 0.50485, saving model to  
model\_save/weights-01-0.5048.hdf5

670/670 [=====] - 8s 9ms/step - loss: 0.6997 -  
accuracy: 0.5026 - auc: 0.5000 - val\_loss: 0.6934 - val\_accuracy: 0.5048 -  
val\_auc: 0.5043 - lr: 0.0100

Epoch 2: LearningRateScheduler setting learning rate to 0.009999999776482582.

Epoch 2/5

668/670 [=====>.] - ETA: 0s - loss: 0.6940 - accuracy:

```
0.5034 - auc: 0.4990
AUC: 0.5 - F1-Score: 0.0
```

```
Epoch 2: val_accuracy improved from 0.50485 to 0.50636, saving model to
model_save/weights-02-0.5064.hdf5
670/670 [=====] - 6s 9ms/step - loss: 0.6940 -
accuracy: 0.5032 - auc: 0.4987 - val_loss: 0.6934 - val_accuracy: 0.5064 -
val_auc: 0.5012 - lr: 0.0100
```

```
Epoch 3: LearningRateScheduler setting learning rate to 0.009499999787658453.
Epoch 3/5
666/670 [=====>.] - ETA: 0s - loss: 0.6941 - accuracy:
0.4953 - auc: 0.4908
AUC: 0.5 - F1-Score: 0.0
```

```
Epoch 3: ReduceLROnPlateau reducing learning rate to 0.0009499999694526196.
```

```
Epoch 3: val_accuracy did not improve from 0.50636
670/670 [=====] - 7s 10ms/step - loss: 0.6941 -
accuracy: 0.4952 - auc: 0.4911 - val_loss: 0.6935 - val_accuracy: 0.4995 -
val_auc: 0.4999 - lr: 9.5000e-04
Epoch 3: early stopping
```

```
[24]: <keras.callbacks.History at 0x7f0ddf3dae10>
```

```
[25]: %tensorboard --logdir logs/fits
```

```
<IPython.core.display.Javascript object>
```

Model-2

1. Use relu as an activation for every layer except output layer.
2. Use SGD with momentum as optimizer.
3. Use RandomUniform(0,1) as initializer.
4. Analyze your output and training process.

```
[26]: %load_ext tensorboard
# Clear any logs from previous runs
!rm -rf ./logs/
```

The tensorboard extension is already loaded. To reload it, use:

```
%reload_ext tensorboard
```

```
[27]: m2_ki = tf.keras.initializers.RandomUniform(minval=0, maxval=1)
```

```
[28]: input_layer = Input(shape=(2,))
layer1 = Dense(units=20, activation='relu',
               ↪kernel_initializer=m2_ki)(input_layer)
layer2 = Dense(units=15, activation='relu', kernel_initializer=m2_ki)(layer1)
layer3 = Dense(units=10, activation='relu', kernel_initializer=m2_ki)(layer2)
```



```

layer4 = Dense(units=5, activation='relu', kernel_initializer=m2_ki)(layer3)
layer5 = Dense(units=2, activation='relu', kernel_initializer=m2_ki)(layer4)
output_layer = Dense(units=1, activation='sigmoid',
    ↪kernel_initializer=m2_ki)(layer5)

```

```

[29]: model2 = Model(inputs=input_layer, outputs=output_layer)
optimizer = tf.keras.optimizers.SGD(learning_rate=0.01, momentum=0.9)
model2.compile(optimizer=optimizer, loss='binary_crossentropy',
    ↪metrics=['accuracy', 'AUC'])

```

```

[30]: score_callback = ScoreCallback(validation_data=(X_test, y_test))

termination_callback = TerminateCallback()

decay_lr_callback = ReduceLROnPlateau(monitor='val_accuracy', factor=0.1,
    ↪patience=1, verbose=1, mode='auto')

lr_schedule_callback = LearningRateScheduler(schedule=scheduler, verbose=1)

filepath = "model_save/weights-{epoch:02d}-{val_accuracy:.4f}.hdf5"
model_save_callback = ModelCheckpoint(filepath=filepath,
    ↪monitor='val_accuracy', verbose=1, save_best_only=True, mode='auto')

early_stop_callback = EarlyStopping(monitor='val_accuracy', min_delta=0.01,
    ↪patience=2, verbose=1)

log_dir = os.path.join('logs', 'fits', datetime.datetime.now().
    ↪strftime("%Y%m%d-%H%M%S"))
tensorboard_callback = tf.keras.callbacks.TensorBoard(log_dir=log_dir,
    ↪histogram_freq=1)

callbacks = [score_callback, termination_callback,
            decay_lr_callback, lr_schedule_callback,
            model_save_callback, early_stop_callback, tensorboard_callback]

```

```

[31]: model2.fit(x=X_train, y=y_train, epochs=5, validation_data=(X_test, y_test),
    ↪batch_size=20, callbacks=callbacks)

```

Epoch 1: LearningRateScheduler setting learning rate to 0.009999999776482582.

Epoch 1/5

1/670 [...] - ETA: 7:46 - loss: 115.2181 -  
accuracy: 0.4000 - auc: 0.4495

WARNING:tensorflow:Callback method `on\_train\_batch\_end` is slow compared to the batch time (batch time: 0.0016s vs `on\_train\_batch\_end` time: 0.0035s). Check your callbacks.

```
649/670 [=====>.] - ETA: 0s - loss: 0.9067 - accuracy:
0.4995 - auc: 0.4969
AUC: 0.5 - F1-Score: 0.0
```

```
Epoch 1: val_accuracy improved from -inf to 0.50636, saving model to
model_save/weights-01-0.5064.hdf5
```

```
670/670 [=====] - 3s 4ms/step - loss: 0.9001 -
accuracy: 0.4992 - auc: 0.4966 - val_loss: 0.6931 - val_accuracy: 0.5064 -
val_auc: 0.5000 - lr: 0.0100
```

```
Epoch 2: LearningRateScheduler setting learning rate to 0.009999999776482582.
Epoch 2/5
```

```
644/670 [=====>..] - ETA: 0s - loss: 0.6935 - accuracy:
0.5031 - auc: 0.4989
AUC: 0.5 - F1-Score: 0.0
```

```
Epoch 2: ReduceLROnPlateau reducing learning rate to 0.0009999999776482583.
```

```
Epoch 2: val_accuracy did not improve from 0.50636
```

```
670/670 [=====] - 3s 4ms/step - loss: 0.6935 -
accuracy: 0.5016 - auc: 0.4981 - val_loss: 0.6936 - val_accuracy: 0.4936 -
val_auc: 0.5000 - lr: 1.0000e-03
```

```
Epoch 3: LearningRateScheduler setting learning rate to 0.000949999934528023.
Epoch 3/5
```

```
667/670 [=====>.] - ETA: 0s - loss: 0.6932 - accuracy:
0.5015 - auc: 0.4992
AUC: 0.5 - F1-Score: 0.0
```

```
Epoch 3: ReduceLROnPlateau reducing learning rate to 9.499999578110874e-05.
```

```
Epoch 3: val_accuracy did not improve from 0.50636
```

```
670/670 [=====] - 3s 4ms/step - loss: 0.6932 -
accuracy: 0.5012 - auc: 0.4991 - val_loss: 0.6932 - val_accuracy: 0.4936 -
val_auc: 0.5000 - lr: 9.5000e-05
```

```
Epoch 3: early stopping
```

```
[31]: <keras.callbacks.History at 0x7f0ddb7c8750>
```

```
[32]: %tensorboard --logdir logs/fits
```

```
Reusing TensorBoard on port 6006 (pid 1993), started 0:26:43 ago. (Use '!kill_
1993' to kill it.)
```

```
<IPython.core.display.Javascript object>
```

```
Model-3
```

1. Use relu as an activation for every layer except output layer.
2. Use SGD with momentum as optimizer.

3. Use `he_uniform()` as initializer.
4. Analyze your output and training process.

```
[33]: %load_ext tensorboard
      # Clear any logs from previous runs
      !rm -rf ./logs/
```

The tensorboard extension is already loaded. To reload it, use:

```
%reload_ext tensorboard
```

```
[34]: m3_ki = tf.keras.initializers.HeUniform()
```

```
[35]: input_layer = Input(shape=(2,))
      layer1 = Dense(units=20, activation='relu',
      ↪kernel_initializer=m3_ki)(input_layer)
      layer2 = Dense(units=15, activation='relu', kernel_initializer=m3_ki)(layer1)
      layer3 = Dense(units=10, activation='relu', kernel_initializer=m3_ki)(layer2)
      layer4 = Dense(units=5, activation='relu', kernel_initializer=m3_ki)(layer3)
      layer5 = Dense(units=2, activation='relu', kernel_initializer=m3_ki)(layer4)
      output_layer = Dense(units=1, activation='sigmoid',
      ↪kernel_initializer=m3_ki)(layer5)
```

```
[36]: model3 = Model(inputs=input_layer, outputs=output_layer)
      optimizer = tf.keras.optimizers.SGD(learning_rate=0.01, momentum=0.9)
      model3.compile(optimizer=optimizer, loss='binary_crossentropy',
      ↪metrics=['accuracy', 'AUC'])
```

```
[37]: score_callback = ScoreCallBack(validation_data=(X_test, y_test))

      termination_callback = TerminateCallBack()

      decay_lr_callback = ReduceLROnPlateau(monitor='val_accuracy', factor=0.1,
      ↪patience=1, verbose=1, mode='auto')

      lr_schedule_callback = LearningRateScheduler(schedule=scheduler, verbose=1)

      filepath = "model_save/weights-{epoch:02d}-{val_accuracy:.4f}.hdf5"
      model_save_callback = ModelCheckpoint(filepath=filepath,
      ↪monitor='val_accuracy', verbose=1, save_best_only=True, mode='auto')

      early_stop_callback = EarlyStopping(monitor='val_accuracy', min_delta=0.01,
      ↪patience=2, verbose=1)

      log_dir = os.path.join('logs', 'fits', datetime.datetime.now().
      ↪strftime("%Y%m%d-%H%M%S"))
      tensorboard_callback = tf.keras.callbacks.TensorBoard(log_dir=log_dir,
      ↪histogram_freq=1)
```

```
callbacks = [score_callback, termination_callback,
             decay_lr_callback, lr_schedule_callback,
             model_save_callback, early_stop_callback, tensorboard_callback]
```

```
[38]: model3.fit(x=X_train, y=y_train, epochs=5, validation_data=(X_test, y_test),  
               ↪batch_size=20, callbacks=callbacks)
```

Epoch 1: LearningRateScheduler setting learning rate to 0.009999999776482582.

Epoch 1/5

1/670 [...] - ETA: 7:36 - loss: 0.7161 - accuracy:  
0.4500 - auc: 0.4050

WARNING:tensorflow:Callback method `on\_train\_batch\_end` is slow compared to the batch time (batch time: 0.0024s vs `on\_train\_batch\_end` time: 0.0033s). Check your callbacks.

644/670 [=====>..] - ETA: 0s - loss: 0.6489 - accuracy:  
0.6180 - auc: 0.6707  
AUC: 0.5 - F1-Score: 0.0

Epoch 1: val\_accuracy improved from -inf to 0.66500, saving model to  
model\_save/weights-01-0.6650.hdf5

670/670 [=====] - 3s 4ms/step - loss: 0.6489 -  
accuracy: 0.6184 - auc: 0.6704 - val\_loss: 0.6103 - val\_accuracy: 0.6650 -  
val\_auc: 0.7307 - lr: 0.0100

Epoch 2: LearningRateScheduler setting learning rate to 0.009999999776482582.

Epoch 2/5

654/670 [=====>..] - ETA: 0s - loss: 0.6128 - accuracy:  
0.6644 - auc: 0.7243  
AUC: 0.5 - F1-Score: 0.0

Epoch 2: val\_accuracy improved from 0.66500 to 0.66970, saving model to  
model\_save/weights-02-0.6697.hdf5

670/670 [=====] - 2s 3ms/step - loss: 0.6128 -  
accuracy: 0.6644 - auc: 0.7243 - val\_loss: 0.6073 - val\_accuracy: 0.6697 -  
val\_auc: 0.7317 - lr: 0.0100

Epoch 3: LearningRateScheduler setting learning rate to 0.009499999787658453.

Epoch 3/5

656/670 [=====>..] - ETA: 0s - loss: 0.6101 - accuracy:  
0.6633 - auc: 0.7290  
AUC: 0.5 - F1-Score: 0.0

Epoch 3: ReduceLROnPlateau reducing learning rate to 0.0009499999694526196.

Epoch 3: val\_accuracy did not improve from 0.66970

670/670 [=====] - 2s 3ms/step - loss: 0.6099 -

```
accuracy: 0.6636 - auc: 0.7294 - val_loss: 0.6098 - val_accuracy: 0.6642 -  
val_auc: 0.7318 - lr: 9.5000e-04  
Epoch 3: early stopping
```

```
[38]: <keras.callbacks.History at 0x7f0ddc7f1890>
```

```
[39]: %tensorboard --logdir logs/fits
```

```
Reusing TensorBoard on port 6006 (pid 1993), started 0:37:00 ago. (Use '!kill_  
↪1993' to kill it.)
```

```
<IPython.core.display.Javascript object>
```

Model-4

1. Try with any values to get better accuracy/f1 score.

```
[40]: %load_ext tensorboard  
# Clear any logs from previous runs  
!rm -rf ./logs/
```

```
The tensorboard extension is already loaded. To reload it, use:  
%reload_ext tensorboard
```

```
[41]: m4_ki = tf.keras.initializers.HeNormal()
```

```
[42]: input_layer = Input(shape=(2,))  
layer1 = Dense(units=20, activation='relu',  
↪kernel_initializer=m4_ki)(input_layer)  
layer2 = Dense(units=15, activation='relu', kernel_initializer=m4_ki)(layer1)  
layer3 = Dense(units=10, activation='relu', kernel_initializer=m4_ki)(layer2)  
layer4 = Dense(units=5, activation='relu', kernel_initializer=m4_ki)(layer3)  
layer5 = Dense(units=2, activation='relu', kernel_initializer=m4_ki)(layer4)  
output_layer = Dense(units=1, activation='sigmoid',  
↪kernel_initializer=m4_ki)(layer5)
```

```
[43]: model4 = Model(inputs=input_layer, outputs=output_layer)  
optimizer = tf.keras.optimizers.Adam()  
model4.compile(optimizer=optimizer, loss='binary_crossentropy',  
↪metrics=['accuracy', 'AUC'])
```

```
[44]: score_callback = ScoreCallBack(validation_data=(X_test, y_test))  
  
termination_callback = TerminateCallBack()  
  
decay_lr_callback = ReduceLROnPlateau(monitor='val_accuracy', factor=0.1,  
↪patience=1, verbose=1, mode='auto')  
  
lr_schedule_callback = LearningRateScheduler(schedule=scheduler, verbose=1)
```

```

filepath = "model_save/weights-{epoch:02d}-{val_accuracy:.4f}.hdf5"
model_save_callback = ModelCheckpoint(filepath=filepath,
    ↪monitor='val_accuracy', verbose=1, save_best_only=True, mode='auto')

early_stop_callback = EarlyStopping(monitor='val_accuracy', min_delta=0.01,
    ↪patience=2, verbose=1)

log_dir = os.path.join('logs', 'fits', datetime.datetime.now().
    ↪strftime("%Y%m%d-%H%M%S"))
tensorboard_callback = tf.keras.callbacks.TensorBoard(log_dir=log_dir,
    ↪histogram_freq=1)

callbacks = [score_callback, termination_callback,
             decay_lr_callback, lr_schedule_callback,
             model_save_callback, early_stop_callback, tensorboard_callback]

```

```

[45]: model4.fit(x=X_train, y=y_train, epochs=5, validation_data=(X_test, y_test),
    ↪batch_size=20, callbacks=callbacks)

```

Epoch 1: LearningRateScheduler setting learning rate to 0.0010000000474974513.

Epoch 1/5

1/670 [...] - ETA: 12:10 - loss: 0.6720 - accuracy:  
0.5000 - auc: 0.6450

WARNING:tensorflow:Callback method `on\_train\_batch\_end` is slow compared to the batch time (batch time: 0.0014s vs `on\_train\_batch\_end` time: 0.0044s). Check your callbacks.

654/670 [=====>.] - ETA: 0s - loss: 0.6732 - accuracy:  
0.5880 - auc: 0.6580  
AUC: 0.5 - F1-Score: 0.0

Epoch 1: val\_accuracy improved from -inf to 0.64303, saving model to  
model\_save/weights-01-0.6430.hdf5

670/670 [=====] - 4s 4ms/step - loss: 0.6724 -  
accuracy: 0.5900 - auc: 0.6614 - val\_loss: 0.6537 - val\_accuracy: 0.6430 -  
val\_auc: 0.6960 - lr: 0.0010

Epoch 2: LearningRateScheduler setting learning rate to 0.0010000000474974513.

Epoch 2/5

654/670 [=====>.] - ETA: 0s - loss: 0.6316 - accuracy:  
0.6632 - auc: 0.7222  
AUC: 0.5 - F1-Score: 0.0

Epoch 2: val\_accuracy improved from 0.64303 to 0.67136, saving model to  
model\_save/weights-02-0.6714.hdf5

670/670 [=====] - 2s 3ms/step - loss: 0.6311 -

```
accuracy: 0.6637 - auc: 0.7227 - val_loss: 0.6215 - val_accuracy: 0.6714 -  
val_auc: 0.7281 - lr: 0.0010
```

```
Epoch 3: LearningRateScheduler setting learning rate to 0.0009500000451225787.  
Epoch 3/5
```

```
660/670 [=====>.] - ETA: 0s - loss: 0.6158 - accuracy:  
0.6739 - auc: 0.7308  
AUC: 0.5 - F1-Score: 0.0
```

```
Epoch 3: ReduceLROnPlateau reducing learning rate to 9.500000160187483e-05.
```

```
Epoch 3: val_accuracy did not improve from 0.67136  
670/670 [=====] - 2s 3ms/step - loss: 0.6156 -  
accuracy: 0.6746 - auc: 0.7311 - val_loss: 0.6165 - val_accuracy: 0.6689 -  
val_auc: 0.7274 - lr: 9.5000e-05
```

```
Epoch 4: LearningRateScheduler setting learning rate to 9.500000305706635e-05.  
Epoch 4/5
```

```
670/670 [=====] - ETA: 0s - loss: 0.6091 - accuracy:  
0.6737 - auc: 0.7377  
AUC: 0.5 - F1-Score: 0.0
```

```
Epoch 4: ReduceLROnPlateau reducing learning rate to 9.500000305706636e-06.
```

```
Epoch 4: val_accuracy did not improve from 0.67136  
670/670 [=====] - 2s 3ms/step - loss: 0.6091 -  
accuracy: 0.6737 - auc: 0.7377 - val_loss: 0.6153 - val_accuracy: 0.6689 -  
val_auc: 0.7267 - lr: 9.5000e-06
```

```
Epoch 4: early stopping
```

```
[45]: <keras.callbacks.History at 0x7f0ddc502cd0>
```

```
[46]: %tensorboard --logdir logs/fits
```

```
Reusing TensorBoard on port 6006 (pid 1993), started 0:47:17 ago. (Use '!kill_  
1993' to kill it.)
```

```
<IPython.core.display.Javascript object>
```

```
End of the file.
```